


The ABC of computational Text Analysis

10: NLP with Python

Alex Flückiger
7 May 2020

Recap last Lecture

- introduce Python 
editor
syntax

Outline

- NLP, getting serious! 🙇
- interactive coding
interrupt, ask, complement
- mini-project

A brick wall covered in white tiles with various letters and words, serving as a background for the title.

Primer on NLP

What is a Word?

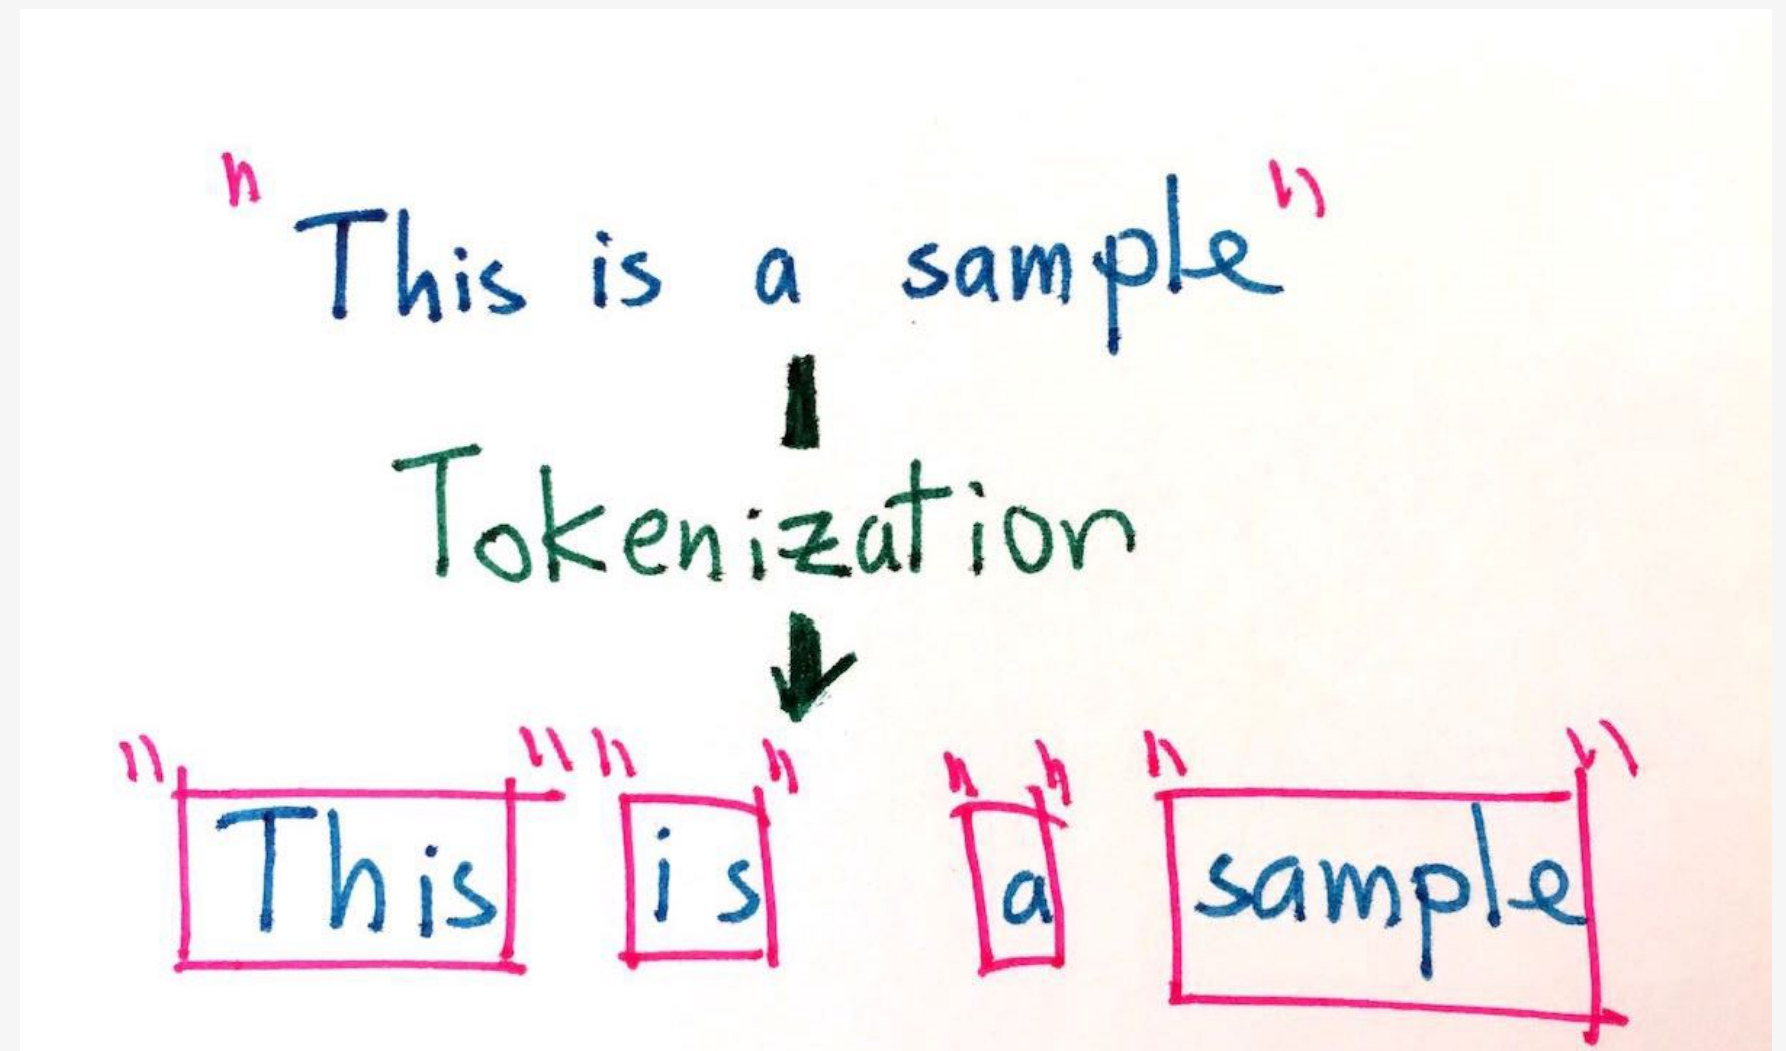
- word ~ segments between whitespace
- yet, there are ...

contractions: *U.S., don't*

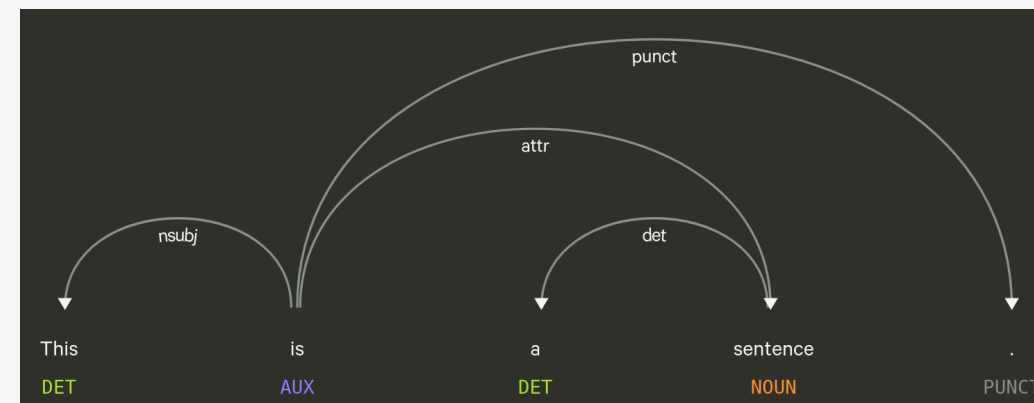
collocations: *New York*

Token

- token ~ computational unit
representation of words
- lemma ~ base form of a word
texts → text
goes → go
- stop words ~ functional words
lack meaning
the, a, on, and



NLP Processing Steps



1. tokenizing

segmenting text into words, punctuations etc.

2. tagging part-of-speech (POS)

assigning word types (e.g. verb, noun)

3. parsing

How does the computer “know?”

- using language specific models
- supervised machine learning
- trained on human-annotated data

Representation of Corpus

document term matrix

- Doc 1: NLP is great. I love NLP.
- Doc 2: I understand NLP.
- Doc 3: NLP, NLP, NLP.

	NLP	I	is	term
	2	1	1	...
	1	1	0	...
	3	0	0	...
Doc ID	term frequency

Mini-Project

- multiple documents
- write a script
- compare ...
historically
across actors
- relative frequency 👍, absolute frequency 👎

Optional Seminar Paper

- writing a seminar paper (4 ECTS)
- get in touch to discuss your idea

Outlook NLP

- explore corpus differences by political party
- term frequency over time
- data

party programmes

1 August speeches by Swiss Federal Councillors

Do NLP

Check out the Python code

Resources

tutorials on spaCy

- [official spaCy 101](#)
- [official online course spaCy](#)
- [Hitchhiker's Guide to NLP in spaCy](#)